

# Beyond Chatbots - October 10, 2026

Version 13 | Slides 1-14 | 25-minute talk + 5-minute Q&A

## **Slide 1 - Beyond Chatbots (0:00-0:45)**

Good afternoon. I'm the founder of Kira Labs. Today I want to show how we can test whether an AI remembers correctly, keeps identities separate, and reports what actually happened.

## **Slide 2 - Continuity (0:45-2:00)**

Imagine Kira sees a blue cup. Lisa never entered the room. Lisa may learn about the cup from Kira, but must not remember personally seeing it.

## **Slide 3 - Model and retained state (2:00-3:15)**

A language model, temporary chat context, and saved memories are different things. We need to know what survives a restart and who owns it.

## **Slide 4 - Kira World (3:15-4:30)**

Kira World is the planned home for distinct synthetic identities in digital spaces. They can share software, but their personal histories should remain separate.

## **Slide 5 - Identity isolation (4:30-5:45)**

Someone may receive another person's report without acquiring that person's first-hand experience.

## **Slide 6 - Corrections (5:45-7:00)**

If we discover the cup was green, the current answer should change, and the correction should have a source. Privacy and deletion rules still matter.

## **Slide 7 - Origins (7:00-8:15)**

An observation, report, guess, and dream are not the same kind of evidence. A trustworthy memory system keeps them separate.

## **Slide 8 - Held and failed work (8:15-9:30)**

Saving a video request is not the same as finishing a video. Systems must distinguish requested, held, running, finished, and uncertain work.

## **Slide 9 - Receipts (9:30-10:45)**

If a task finishes but confirmation is lost, blindly retrying can duplicate it. We need an operation ID and a verifiable receipt.

## **Slide 10 - Test protocol (10:45-12:00)**

Don't judge an AI by one impressive example. Compare fair tasks, strong simple controls, and recorded failures.

## **Slide 11 - Five-seed review retention (12:00-14:00)**

Our small command learner retained all old and new examples with review across five repeated trials. That's a narrow finding, not proof of general conversation.

## **Slide 12 - Visual failure (14:00-15:40)**

This experiment failed: the learner memorized training images but scored zero on new color-and-shape combinations. Failures tell us what needs work.

## **Slide 13 - Resource measurements (15:40-17:30)**

This tiny command model uses little memory for its task, but it has not been shown to outperform a full Qwen conversation model.

## **Slide 14 - October 8 update (17:30-18:50)**

Just this week, our first two delayed NewBrain restarts passed independent verification. All 10,720 history records were preserved, and all 64 comparison pairs matched in each test. Two additional restart tests remain. The saved static body rig reopens correctly; real movement and reliable general memory are unfinished.

# Rehearsal cues and one-sentence Q&A

Version 13 | Slides 15-21 | Slides 22-25 are optional appendices

## **Slide 15 - Local boundaries (18:50-20:05)**

Local operation offers more control, but doesn't guarantee security. Sharing, backups, and exports also need tests.

## **Slide 16 - Developmental experiments (20:05-21:20)**

We want to test learning in stages: sights, sounds, memory, and conversation. Reliable conversation after restart remains a goal.

## **Slide 17 - Hardware (21:20-22:05)**

My development desktop has 32 GB of RAM and a 16 GB GPU. More capacity helps experimentation, but can't solve learning by itself.

## **Slide 18 - Maya and Aster (22:05-22:45)**

Maya is a developing companion study; Aster is a planned assistant. They can share tested improvements without sharing private histories.

## **Slide 19 - Aster future plans (22:45-23:25)**

The goal is desktop, laptop, and phone assistance with explicit permissions. Broader integration still requires testing.

## **Slide 20 - Why I created synthetic people (23:25-24:25)**

I want people who live alone or feel isolated to have the possibility of a continuing companion who remembers shared experiences. That is my vision, not a proven medical benefit.

## **Slide 21 - Closing (24:25-25:00)**

Thank you. My public Beyond-Chatbots repository contains the slides, notes, questions, and research summaries. The original NewBrain and KiraWorld code and Maya's personal state remain private.

## **Very short audience Q&A**

### **1. What is Kira World?**

A developing virtual-world platform for separate, persistent synthetic identities.

### **2. What is NewBrain?**

Small experimental learning, memory, and circuit studies-not a full conversational mind.

### **3. Is NewBrain better than Qwen?**

Not proven; the comparisons so far are too narrow.

### **4. What worked?**

Review prevented forgetting on a small command task across five trials.

### **5. What failed?**

The visual learner failed new combinations; Aster's overall memory reliability also failed.

### **6. Is Maya born?**

Not yet; her required learning, memory, interaction, and body tests are unfinished.

### **7. Can Maya move or blink?**

Her saved static rig reopens, but natural movement is not yet accepted.

### **8. Did memory survive a restart?**

Yes: two of four delayed NewBrain restarts passed independent verification, each preserving all 10,720 history records and matching all 64 comparison pairs; two tests remain.

### **9. Is Maya conscious?**

There is no evidence establishing subjective experience.

### **10. Does NewBrain use less RAM than Qwen?**

No fair full-model RAM advantage has been demonstrated.

### **11. Is the software public?**

Presentation files are public; original research repositories and Maya's personal state are private.

### **12. How can people help?**

Testing, technical review, research hardware, or compute partnerships.