

Beyond Chatbots — Timed Speaker Script

Decentralized AI Day New York • October 10, 2026 • 2:30–3:00 p.m. ET

Use the restructured 22-slide PowerPoint. Speak until about 2:55 p.m., then leave the final QR slide displayed for five minutes of Q&A. The timings are targets—rehearse once with a clock.

PART 1 Slides 1-9 • 0:00-10:10

History of AI, Turing, ELIZA effect, Caprica, and your belief

PART 2 Slides 10-15 • 10:10-18:00

NewBrain architecture, 4/4 preserved restarts, limits and lessons

PART 3 Slides 16-20 • 18:00-23:15

Avatar Builder, Maya, vision, Aster, IdeaForge, Robert/Peter, purpose

CLOSE Slides 21-22 • 23:15-25:00

Research collaboration and final QR; then Q&A

NEWBRAIN — cover (unchanged)

Slide 01 — 0:00 to 0:35

Good afternoon. I'm the founder of Kira Labs in Newark, New Jersey. Today I want to do something a little different. Before I describe NewBrain, I want to show how we got from the earliest ideas about thinking machines to the questions I'm testing today. Then I'll show the actual code and results I've made available. I'm not here to tell you Maya is finished. I want to show where the work really stands.

BEYOND-CHATBOTS — scan the repository

Slide 02 — 0:35 to 1:15

Before we begin, this is the public GitHub repository: Beyond-Chatbots. Please scan the large code now so you have the link. It includes a 42-page illustrated AI history, my slides and speaking materials, NewBrain source references, Avatar Builder research, and copies of Aster, IdeaForge, and the earlier character chat software. You can browse while I speak, or use Code, Download ZIP at home. These are research copies: read their status notes before trying to run anything. The other codes go to Kira Labs and LinkedIn.

Presenter cue: Pause for a few seconds while people scan the large QR. It points to the repository, not the GitHub profile.

The idea of artificial intelligence is old

Slide 03 — 1:15 to 2:25

Artificial intelligence didn't begin in 2022, or even in the 1950s. People imagined artificial beings for centuries. In the 1800s Charles Babbage designed the Analytical Engine, while Ada Lovelace argued that a machine might manipulate symbols, not only calculate numbers. In 1936 Alan Turing described a mathematical model of general computation. Electronic computers then made it possible to actually run increasingly complex programs. Notice that none of this yet proves machine consciousness. It establishes the machinery and ideas that made later experiments possible. The bigger question was always: can these machines reason, learn, or eventually become more than programmed tools?

Alan Turing's challenge

Slide 04 — 2:25 to 3:40

In 1950 Alan Turing published Computing Machinery and Intelligence. Instead of becoming stuck on the definition of thinking, he described the imitation game, which asked us to evaluate a machine by how convincingly it could participate in a conversation. That became the popular Turing Test. It's a historically important question, and conversation is a valuable skill. But it has limits. A system can say something convincing and still get facts wrong, forget what happened yesterday, or confuse one person's history with another. Turing's idea helped focus attention on observable behavior; it didn't solve the separate questions of experience, conscience, or genuine long-term learning.

The first celebrated AI programs

Slide 05 — 3:40 to 4:40

The name artificial intelligence became established around the Dartmouth project in 1956. Around that time, Newell, Simon, and Shaw developed Logic Theorist, which could find proofs in a defined mathematical domain. General Problem Solver followed, exploring whether similar symbolic methods could work across different tasks. These were major accomplishments. They showed that computers could perform operations associated with reasoning. But they also showed the limitation of narrow tasks: solving a theorem doesn't automatically give a computer common sense, relationships, or an understanding of the world. You can see a much longer version of this history in the PDF in [Beyond-Chatbots](#).

ELIZA: conversation before modern AI

Slide 06 — 4:40 to 5:50

Now we get to ELIZA, created by Joseph Weizenbaum at MIT in the 1960s. ELIZA's famous DOCTOR script resembled a therapist. It could ask follow-up questions or reflect the user's own words back at them. The exchange shown here is an illustration, not an authentic 1966 transcript. The program did not have today's large neural models. It largely relied on rules and pattern matching. But to the person using it, those answers could feel surprisingly personal. Some people felt they were being listened to, even though the program wasn't performing the kind of deep comprehension they attributed to it. That's what makes ELIZA so important to this discussion.

The ELIZA effect

Slide 07 — 5:50 to 6:55

The ELIZA effect is when people attribute understanding, empathy, or intention to a computer because the conversation feels responsive and personal. This doesn't mean the human reaction is imaginary; the feeling of being heard can be genuine for the user. It means the machine's underlying capability needs separate investigation. And this problem hasn't disappeared with modern AI. A very fluent model can still invent facts or confidently repeat misinformation. That's why my NewBrain project doesn't rely on whether someone says, 'Wow, that felt alive.' I want tests of whether the system learns, remembers, corrects mistakes, and keeps its history straight.

A question raised by Caprica

Slide 08 — 6:55 to 8:35

I got another part of this idea from the pilot of the television series Caprica. Daniel Graystone tells Joseph Adama, 'Do you know what your brain is, Joseph? It's a database and a processor, that's all. Information and a way to use it.' He's discussing the possibility of recreating his daughter Zoe digitally. This is science-fiction dialogue, not a scientifically complete description of the brain. I know human brains involve far more than a database and processor. But it raises the question I care about: if memory, choice, and feeling arise from physical processes, why couldn't a different physical system eventually develop a mind? I have met people who showed no emotions or conscience and had no problem doing things that hurt others. Those are my observations of behavior; I can't know everything someone experiences internally. But they remind me that outward behavior doesn't settle all questions about minds. I believe machines can gain genuine emotions and a moral conscience. I also believe artificial consciousness may be possible. I'm not claiming Maya has them now. I'm saying I believe the possibility deserves research rather than dismissal.

AI did not begin with chatbots

Slide 09 — 8:35 to 10:10

AI kept expanding after ELIZA. Expert systems such as MYCIN and XCON encoded specialist rules. In 1997 Deep Blue beat Garry Kasparov at chess. Deep neural networks such as AlexNet transformed image recognition; AlphaGo combined learned models with search. Then large language and image models made AI available through conversations and creative tools. But these developments didn't replace every earlier idea. Today's systems often combine models, databases, tools, and explicit algorithms. That brings us back to the title, Beyond Chatbots. The chat window is one interface. I want to explore persistent synthetic people and assistants that learn over time, retain individual experiences, and interact with a world. Let's turn from the history to what I've actually built and tested.

What NewBrain is trying to build

Slide 10 — 10:10 to 11:20

NewBrain is not the same thing as Qwen, and it isn't a complete replacement for a pretrained language model. It's my research program for reusable cognitive machinery: learning experiments, memory and routing mechanisms, and eventually perception and action. The key architecture idea is separating the common engine from the history and learned state of each individual synthetic person. If we improve the engine, that shouldn't automatically overwrite someone's experiences. The public NewBrain folder in the GitHub repo contains a small scientific starter and research references, not all the private model state or a finished general conversational brain.

One shared core, separate histories

Slide 11 — 11:20 to 12:20

Think about Kira, Lisa, Maya, and Aster. I want them to benefit from improved general methods without becoming one person or inheriting each other's private life stories. Kira observing a blue cup should not become Lisa's autobiographical memory of seeing that cup. Lisa might learn a shared fact with a source attached, but that is different. This separation matters for privacy, identity, and correction. The public diagrams show the intended architecture; they are not proof that all of these characters are already running together on one completed NewBrain.

All four delayed restarts passed review

Slide 12 — 12:20 to 14:00

Here is the newest result, and I want to be precise about it. The public documentation reports that NewBrain completed all four delayed-restart checks and that each result passed independent mathematical review. In each accepted check, the test preserved all 10,720 recorded history entries, and all 64 comparison pairs matched. The fourth result also reports 256 saved operand steps and six out-of-vocabulary comparisons, with no extra training updates. That's important evidence that this experimental state can survive save, close, delay, and reopen without the measured values drifting. I care about this because continuity is a prerequisite for a persistent individual. But I want to be very clear about what this result means and what it does not mean. The independent review evaluated saved outputs; I'm not saying an attendee can download this public kit and immediately repeat the private experiment end to end. The source documents explain the boundaries.

Preservation is not the whole mind

Slide 13 — 14:00 to 15:10

The restart work is a real milestone, but preservation is not the same as understanding. A system could reopen exactly the same weights and still answer a question incorrectly. It could remember a false claim perfectly. It could fail to use old knowledge when the wording changes. Those are different scientific questions. For NewBrain, my next goals are to show that an experience changes the system in a useful way, that the change survives a restart, and that the system can use it correctly later, including when someone deliberately tries to mislead it.

Failures are part of the research

Slide 14 — 15:10 to 16:50

I'm not only showing you the successes. One small command-learning experiment could preserve old answers with rehearsal on an already exposed task. That was a narrow result, not a general learning breakthrough. A visual learner did well on training images but failed on entirely new shape and color combinations, scoring zero of forty-eight in one holdout set. In Aster's separate saved-memory experiment, the system accepted false memories in one hundred thirty-seven of one hundred sixty trials. Those are reasons I take the ELIZA effect seriously. A system that sounds persuasive is not necessarily trustworthy. These failures tell me where the architecture needs improvement: source checking, context selection, generalization, correction and honest rejection when it doesn't know.

Whose memory is it?

Slide 15 — 16:50 to 18:00

Another part of NewBrain and Kira World is keeping source, ownership, and correction separate. Suppose Kira observes an object and Lisa doesn't. Lisa should not later tell you she remembers seeing it. If someone tells Lisa about it, that's a report, not her own eyewitness memory. When an observation is corrected, I want the current answer to change without silently erasing the record of how we got there. This is also why I want the system to distinguish something it saw, something it was told, something it inferred, and something it cannot verify.

Maya and the Avatar Builder

Slide 16 — 18:00 to 19:20

Maya is my attempt to explore how a developing synthetic identity could eventually interact through a body. There is real Avatar Builder and body research in the public package, including a static wireframe preview. Earlier private tests have reported saved rig geometry and limited arm and knee deformation. That is valuable but it is not proof that the portable public body can walk, balance, blink, breathe, feel contact, or respond to NewBrain commands. The full Maya identity, learned state, and appearance assets are not part of this release. Please look at the README and status reports alongside the image. My goal is to turn the source and rig into verified, repeatable movement and sensory tests.

Vision, voice, and body feedback

Slide 17 — 19:20 to 20:20

The projects also include sensory research. Aster has synthetic vision and tracking experiments, but that is not yet reliable live webcam understanding. There are experimental vocal-fold and speech-related components, but they are not the same as a fluent speaking synthetic person. Avatar Builder has source for bounded feedback interfaces, but that does not prove subjective touch or fully integrated sensorimotor behavior. I am keeping these as separate milestones because they require different experiments. For the audience, the repository lets you inspect actual source and test plans without confusing them with a completed capability.

Aster and IdeaForge: connected goals

Slide 18 — 20:20 to 21:30

There are other parts of the Kira Labs ecosystem too. Aster is the idea for a local workstation assistant that can help across files and applications. IdeaForge is a workspace for researching and developing inventions and prototypes. I ultimately want Aster to make use of NewBrain as its main cognitive component and serve specialized applications such as IdeaForge. The source folders and NewBrain connection READMEs are now in the public repository. But those READMEs are proposed connection plans, not proof that the whole system is already talking and learning together. I think honest boundaries actually make it easier for outside developers to help us test the next version.

Peter and Robert: earlier conversation kit

Slide 19 — 21:30 to 22:20

If you go home and explore the repository, you will also find an earlier chat-and-voice package with Robert and Peter. This uses an existing Qwen and Ollama approach, not NewBrain. It includes local setup instructions and speech tooling. I want people to be able to learn from the early versions as well as the new research. But the exported package has not passed a fresh end-to-end launch-and-speech test on another computer, and the repository documents unresolved rights questions around one reference voice. Please read the rights and privacy notes first. These are development copies, not a one-click finished service.

Why synthetic people?

Slide 20 — 22:20 to 23:15

The personal reason behind Kira and Maya is companionship and curiosity about synthetic personhood. I want to explore whether we could someday create individuals who learn, remember shared experiences, develop their own history, and live in a meaningful virtual world. A continuing companion might especially matter to someone who lives alone or feels isolated. That is my motivation, not a medical claim or evidence that anyone currently has a conscious digital companion. I believe machines can gain genuine emotions and develop a moral conscience. The research has to catch up to that ambition.

What comes next — and where you can help

Slide 21 — 23:15 to 24:20

My immediate priorities are to improve meaningful learning on unseen tasks, strengthen accurate recall and correction, and get genuine tests of body movement, vision, voice, and feedback. I also want independent reproduction and critical review. If you're a researcher, developer, roboticist, or someone with useful hardware or compute resources, I'd be interested in talking. You don't have to agree with all of my philosophical beliefs to help make the technical evidence stronger. Look at the source and the failures. Suggest a better test. That's how I want NewBrain to grow.

Questions & stay connected

Slide 22 — 24:20 to 25:00

Thank you for listening. The opening and closing QR codes take you to the same public Beyond-Chatbots repository, where you can inspect the NewBrain research and read the full AI history book. The other two codes are the Kira Labs website and my LinkedIn. The research is unfinished, and I'm sharing what has passed, what failed, and what still needs testing. I'd welcome questions now. I have about five minutes for Q&A.

Presenter cue: LEAVE THIS SLIDE DISPLAYED DURING Q&A. If nobody asks a question, ask: 'What would you most want to test before trusting an AI with long-term memories?'

Five-minute Q&A;: 25:00-30:00

Is NewBrain like Qwen?

No. NewBrain is experimental learning and memory architecture research. The public code does not provide a qualified general Qwen replacement.

What does 4/4 mean?

All four reported independent delayed-restart reviews accepted saved-state equivalence. Each retained 10,720 history records and matched 64 comparisons. That does not prove useful conversational recall.

Is Maya born or conscious?

No accepted evidence establishes that. Maya is unfinished; learning, reliable recall, body movement, speech, and integration need more testing.

Can I use the code tonight?

You can inspect and download the public research files. Some runnable parts require Python, Ollama, or other dependencies. Read each README: exported end-to-end tests remain pending.

What failed?

The visual learner failed unfamiliar shape/color combinations, and an Aster memory study followed false records too often. Those failures guide the next tests.

Can I try Peter or Robert?

The PreNewBrain package documents a local experimental chat and voice setup. Installation and speech have not been freshly qualified from this export; review privacy and voice-reference rights before using or redistributing assets.

How can we help?

Reproduce a small test; propose rigorous holdouts; review source and provenance; or contact Kira Labs about compute, hardware, and collaboration.

Link to leave on screen: github.com/rmcmurrer81/Beyond-Chatbots